

solution

December 19, 2017

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
from scipy import stats
import math
import graphlab
%matplotlib inline
plt.rcParams["figure.figsize"] = [12,9]
```

RuntimeError

Traceback (most recent call last)

RuntimeError: module compiled against API version 0xb but this version of numpy is 0xa

1 task

- Load house sales data: kc_house_data.csv.zip
- What is the content, could you read it? do you understand columns?
- Explore the data for housing
- make scatter plot of selected features
- create simple regression model of sqft_living to price
- evaluate a simple model
- is linear function good enough? try quadratic polynomial

Read in the data. Just to check and know all the columns we visualize the SFrame. To later be able to add a constant part to the fits, we are adding a constant valued column. and for the square fit we create a column with the squared values of 'sqft_living'.

```
In [2]: houses = graphlab.SFrame("../lectures/data/kc_house_data.csv")
houses['const'] = np.ones(len(houses['id']))
houses['sqft_living_sq'] = houses['sqft_living']**2
houses
```

This non-commercial license of GraphLab Create for academic use is assigned to niklas.eicker@rwth

```
[INFO] graphlab.cython.cy_server: GraphLab Create v2.1 started. Logging: /tmp/graphlab_server_15
```

```
Finished parsing file /home/nikl/DataScience/lectures/data/kc_house_data.csv
```

```
Parsing completed. Parsed 100 lines in 0.257914 secs.
```

```
-----  
Inferred types from first 100 line(s) of file as  
column_type_hints=[int,str,float,int,float,int,int,float,int,int,int,int,int,int,int,int,int,flo  
If parsing fails due to incorrect types, you can correct  
the inferred type list above and pass it to read_csv in  
the column_type_hints argument  
-----
```

```
Finished parsing file /home/nikl/DataScience/lectures/data/kc_house_data.csv
```

```
Parsing completed. Parsed 21613 lines in 0.263838 secs.
```

```
Out[2]: Columns:
```

id	int
date	str
price	float
bedrooms	int
bathrooms	float
sqft_living	int
sqft_lot	int
floors	float
waterfront	int
view	int
condition	int
grade	int
sqft_above	int
sqft_basement	int
yr_built	int
yr_renovated	int
zipcode	int
lat	float
long	float
sqft_living15	int
sqft_lot15	int
const	float
sqft_living_sq	float

Rows: 21613

Data:

id		date		price		bedrooms		bathrooms		sqft_living					
7129300520	20141013T000000	221900.0	3	1.0	1180	6414100192	20141209T000000	538000.0	3	2.25	2570				
5631500400	20150225T000000	180000.0	2	1.0	770	2487200875	20141209T000000	604000.0	4	3.0	1960				
1954400510	20150218T000000	510000.0	3	2.0	1680	7237550310	20140512T000000	1225000.0	4	4.5	5420				
1321400060	20140627T000000	257500.0	3	2.25	1715	2008000270	20150115T000000	291850.0	3	1.5	1060				
2414600126	20150415T000000	229500.0	3	1.0	1780	3793500160	20150312T000000	323000.0	3	2.5	1890				
sqft_lot		floors		waterfront		view		condition		grade		sqft_above		sqft_basement	
5650	1.0	0	0	3	7	1180	0	7242	2.0	0	0	3	7	2170	400
10000	1.0	0	0	3	6	770	0	5000	1.0	0	0	5	7	1050	910
8080	1.0	0	0	3	8	1680	0	101930	1.0	0	0	3	11	3890	1530
6819	2.0	0	0	3	7	1715	0	9711	1.0	0	0	3	7	1060	0
7470	1.0	0	0	3	7	1050	730	6560	2.0	0	0	3	7	1890	0
yr_built		yr_renovated		zipcode		lat		long		sqft_living15		...			
1955	0	98178	47.5112	-122.257	1340	...	1951	1991	98125	47.721	-122.319	1690	...		
1933	0	98028	47.7379	-122.233	2720	...	1965	0	98136	47.5208	-122.393	1360	...		
1987	0	98074	47.6168	-122.045	1800	...	2001	0	98053	47.6561	-122.005	4760	...		
1995	0	98003	47.3097	-122.327	2238	...	1963	0	98198	47.4095	-122.315	1650	...		
1960	0	98146	47.5123	-122.337	1780	...	2003	0	98038	47.3684	-122.031	2390	...		

[21613 rows x 23 columns]

Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.

Create a scatter plot of house price and sqft_living.

```
In [3]: graphlab.canvas.set_target('ipynb')
        houses.show(view="Scatter Plot",x="sqft_living",y="price")
```

We will do the linear and quadratic fit on one feature (sqft_living) with graphlab.

```
In [4]: sqft_model_lin = graphlab.linear_regression.create(houses, validation_set=None, target='price')
```

Linear regression:

Number of examples : 21613

Number of features : 1

Number of unpacked features : 1

Number of coefficients : 2

Starting Newton Method

+-----+-----+-----+-----+-----+

Iteration	Passes	Elapsed Time	Training-max_error	Training-rmse	
-----------	--------	--------------	--------------------	---------------	--

+-----+-----+-----+-----+-----+

1	2	1.063856	4362074.683616	261440.790302	
---	---	----------	----------------	---------------	--

+-----+-----+-----+-----+-----+

SUCCESS: Optimal solution found.

Now a model that takes the squared value of 'sqft_living' into account.

```
In [5]: sqft_model_quad = graphlab.linear_regression.create(houses, validation_set=None, target=
```

Linear regression:

Number of examples : 21613

Number of features : 2

Number of unpacked features : 2

Number of coefficients : 3

Starting Newton Method

+-----+-----+-----+-----+-----+

Iteration	Passes	Elapsed Time	Training-max_error	Training-rmse	
-----------	--------	--------------	--------------------	---------------	--

+-----+-----+-----+-----+-----+

1	2	0.075018	5913021.143248	250948.367620	
---	---	----------	----------------	---------------	--

+-----+-----+-----+-----+-----+

SUCCESS: Optimal solution found.

```
In [6]: # RMS of the lin fit
        print sqft_model_lin.evaluate(houses)
```

```
{'max_error': 4362074.683615588, 'rmse': 261440.79030169296}
```

```
In [7]: # RMS of the quadratic fit
        print sqft_model_quad.evaluate(houses)
```

```
{'max_error': 5913021.143247618, 'rmse': 250948.36761971583}
```

RMS of the quadratic method is a little bit better than the linear one, but still very big compared to the values.

2 task

- Split your data into training sample and test sample
- what is training error and testing error of your model?
- predict the house price for a given sqft_living
- predict the sqft_living for a given price of the house

```
In [8]: training, testing = houses.random_split(.8, seed=0)
```

recalculate our linear model on the test data

```
In [9]: sqft_model_lin = graphlab.linear_regression.create(training, validation_set=None, target=
```

Linear regression:

Number of examples : 17384

Number of features : 1

Number of unpacked features : 1

Number of coefficients : 2

Starting Newton Method

```

-----

+-----+-----+-----+-----+-----+
| Iteration | Passes   | Elapsed Time | Training-max_error | Training-rmse |
+-----+-----+-----+-----+-----+
| 1         | 2        | 0.045954     | 4349521.915863     | 262943.613519 |
+-----+-----+-----+-----+-----+

```

SUCCESS: Optimal solution found.

Now compare the rms on training data with the rms on test data

```

In [10]: print "training: " + str(sqft_model_lin.evaluate(training)['rmse'])
         print "testing: " + str(sqft_model_lin.evaluate(testing)['rmse'])

```

```

training: 262943.613519
testing: 255191.027487

```

The errors are comparable...

2.1 Methods to predict prices and sqft_living with the fit parameters.

roots helps solving for the sqft_living values

```

In [11]: def get_house_price(sqft_living):
         return (sqft_model_lin.get('coefficients')['value'][0] + sqft_model_lin.get('coefficients')['value'][1]*sqft_living)

In [12]: def get_house_sqft(price):
         coeff = [sqft_model_lin.get('coefficients')['value'][1], (sqft_model_lin.get('coefficients')['value'][0] - price)]
         return np.roots(coeff)[0]

In [13]: print "predict house sqft_living for prices:"
         print "    500.000 : " + str(get_house_sqft(500000))
         print "    1.000.000 : " + str(get_house_sqft(1000000))
         print "    1.500.000 : " + str(get_house_sqft(1500000))
         print "    2.500.000 : " + str(get_house_sqft(2500000))

```

```
predict house sqft_living for prices:
```

```
500.000 : 1940.41067055
1.000.000 : 3713.72536216
1.500.000 : 5487.04005378
2.500.000 : 9033.66943701
```

```
In [14]: print "predict house prices for given sqft:"
        print "    1000 : "+ str(get_house_price(1000))
        print "    2000 : "+ str(get_house_price(2000))
        print "    3000 : "+ str(get_house_price(3000))
        print "    4000 : "+ str(get_house_price(4000))
```

```
predict house prices for given sqft:
```

```
1000 : 234843.82806
2000 : 516801.679289
3000 : 798759.530518
4000 : 1080717.38175
```

3 task

- add more features
- is the model better now?
- maybe using range of data would work better?

```
In [15]: sqft_model = graphlab.linear_regression.create(training, target='price', features=['sqft_living', 'sqft_basement'],
               sqft_model.evaluate(training))
```

PROGRESS: Creating a validation set from 5 percent of training data. This may take a while.
You can set ``validation_set=None`` to disable validation tracking.

```
Linear regression:
```

```
-----
```

```
Number of examples          : 16450
```

```
Number of features          : 4
```

```
Number of unpacked features : 4
```


Number of coefficients : 5

Starting Newton Method

+-----+-----+-----+-----+-----+-----+					
Iteration	Passes	Elapsed Time	Training-max_error	Validation-max_error	Training-rms
+-----+-----+-----+-----+-----+-----+					
1	2	0.106160	4502064.109799	2987342.356958	231173.72689
+-----+-----+-----+-----+-----+-----+					

SUCCESS: Optimal solution found.

Out[15]: {'max_error': 4502064.109799153, 'rmse': 232189.4072274167}

Yes more featerues lowered our error, but it is still quite high.

```
In [16]: range_houses = houses[(houses['sqft_living'] <= 5000) & (houses['price'] <= 2000000)]
         range_training , range_testing = range_houses.random_split(.8,seed=0)
         sqft_model_selected = graphlab.linear_regression.create(range_training, target='price',
         sqft_model_selected.evaluate(range_training)
```

PROGRESS: Creating a validation set from 5 percent of training data. This may take a while.
You can set ``validation_set=None`` to disable validation tracking.

Linear regression:

Number of examples : 16302

Number of features : 4

Number of unpacked features : 4

Number of coefficients : 5

Starting Newton Method

```
-----  
+-----+-----+-----+-----+-----+-----+  
| Iteration | Passes   | Elapsed Time | Training-max_error | Validation-max_error | Training-rms  
+-----+-----+-----+-----+-----+-----+  
| 1         | 2        | 0.104579     | 1330911.700288     | 983736.696828       | 179574.56947  
+-----+-----+-----+-----+-----+-----+  
  
SUCCESS: Optimal solution found.
```

```
Out[16]: {'max_error': 1330911.7002882427, 'rmse': 179424.43561122433}
```

on our limited range for prices < 2mil and sqft_living < 5000 the model works much better! On the other hand the points are now also closer together anyways.

4 task

- predict house price for a house id = 5309101299 (does not exists! using 2008000270 instead)
- what is this house like?
- predict house price for a house id = 1925069082

```
In [17]: houses[houses['id'] == 2008000270]
```

```
Out[17]: Columns:
```

```
id          int
date        str
price       float
bedrooms    int
bathrooms   float
sqft_living int
sqft_lot    int
floors      float
waterfront  int
view        int
condition   int
grade       int
sqft_above  int
sqft_basement int
yr_built    int
yr_renovated int
zipcode     int
lat         float
long        float
sqft_living15 int
sqft_lot15  int
const       float
sqft_living_sq float
```

```
Rows: Unknown
```

```
Data:
```

```
+-----+-----+-----+-----+-----+-----+
| id | date | price | bedrooms | bathrooms | sqft_living |
+-----+-----+-----+-----+-----+-----+
| 2008000270 | 20150115T000000 | 291850.0 | 3 | 1.5 | 1060 |
+-----+-----+-----+-----+-----+-----+
| sqft_lot | floors | waterfront | view | condition | grade | sqft_above | sqft_basemen
+-----+-----+-----+-----+-----+-----+-----+
| 9711 | 1.0 | 0 | 0 | 3 | 7 | 1060 | 0
+-----+-----+-----+-----+-----+-----+-----+
| yr_built | yr_renovated | zipcode | lat | long | sqft_living15 | ... |
```

```
+-----+-----+-----+-----+-----+-----+
| 1963 |      0      | 98198 | 47.4095 | -122.315 |      1650      | ... |
+-----+-----+-----+-----+-----+-----+
[? rows x 23 columns]
Note: Only the head of the SFrame is printed. This SFrame is lazily evaluated.
You can use sf.materialize() to force materialization.
```

4.1 house nr 2008000270

It is a rather small house, with 1060 sqft living and only one floor. The price is one of the lowest of the dataset with 291850 dollar. There are 3 bedrooms and 1.5(?) bathrooms. This house was build in 1963.

```
In [18]: print "predict price of house nr. 1925069082"
         print "          predicted: "+ str(get_house_price(houses[houses['id'] == 1925069082]['s
         print "          real: "+ str(houses[houses['id'] == 1925069082]['price'][0])
```

```
predict price of house nr. 1925069082
predicted: 1261170.40653
real: 2200000.0
```

The prediciton for the price is only ~6% off